

Scaling up

Before we load test, add some resource limits to your pods. We're doing this to get around this error - <https://github.com/kubernetes-sigs/metrics-server/issues/237>

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: cpu-deployment
spec:
  replicas: 2
  selector:
    matchLabels:
      app: cpu-app
  template:
    metadata:
      labels:
        app: cpu-app
    spec:
      containers:
        - name: cpu-app
          image: 100xdevs/week-28:latest
          ports:
            - containerPort: 3000
          resources:
            requests:
              cpu: "100m"
            limits:
              cpu: "1000m"
```



- Try sending a bunch of requests to the server (or just visit it in the browser)

```
nom i -a loadtest
```

↗

<

>

↑

0



- Check the CPU usage

```
kubectl top pods
```



- See the hpa average usage

```
kubectl get hpa
```



- Check the number of pods, see as they scale up

```
kubectl get pods
```

