

Kubernetes Lab

Base repository - <https://github.com/code100x/algorithmic-arena/>

Things to do -

1. Create a PV, PVC for the postgres database.
2. Create a PV, PVC for redis.
3. Create deployments for redis, postgres.
4. Create ClusterIP services for redis, postgres
5. Create a deployment for the nextjs app, expose it via a loadbalancer service on ▶
6. Create a deployment for the judge api server. Expose it via a ClusterIP service
7. Create a deployment for the judge workers. Add resource limits and requests to it
8. Create a HPA that scales based on the **pending submission queue** length in the redis queue
 1. You can either expose an endpoint that you use as a custom metric
 2. You can put all metrics in prometheus and pick them up from there
 3. You can use KEDA to scale based on redis queue length



